

Using Benchmark Datasets to Correct for Measurement Error in Privacy-Preserving Classifiers of AI Usage

Jacob Carlson
Harvard University*

July 2026

1 Introduction

Several recent working papers from the economics teams at leading AI labs are interested in measuring economically salient attributes of AI usage data. Typically, they do so using automated, *privacy-preserving* classifiers of AI usage/session data, i.e., frontier LLMs prompted to produce classifications or estimates of quantities of interest based on the session text transcripts and other metadata.

Example 1 (Measuring coding task complexity [Johnston et al., 2026]). Johnston et al. [2026] introduce a privacy-preserving classifier (gpt-5-mini) of Codex sessions to measure how “complex” a user request submitted to Codex is. The prompt solicits the “active work time in minutes an experienced human would need” to complete the task, as a specific quantitative characterization of task complexity. $\triangle$

Example 2 (Measuring task value [Hitzig et al., 2026]). Hitzig et al. [2026] format a Claude code session’s core task as a freelance job posting (using Claude Haiku 4.5) and then use a privacy-preserving classifier (Claude Opus 4.7) to estimate a price for the session-as-posting. (This is a few-shot classification, in which the prompt for classification is supplemented with a “calibration corpus” of 200 real freelance postings with their actual prices.) $\triangle$

Example 3 (Measuring coding session success [Hitzig et al., 2026]). Hitzig et al. [2026] use an ensemble of privacy-preserving classifiers (Claude Sonnet 4.6 with varying classification prompts) to measure whether or not a Claude code session ended as a success (or partial success, or failure, or was abandoned, etc.). $\triangle$

Of course, in any of these examples, what the classifier outputs is a prediction, not the “ground truth” measurement of interest. As such, there is a measurement error question at the heart of analyses that are downstream of classifier-based measurements: if for some reason the privacy-preserving classifier is systematically biased when making predictions (even if it is

*Email: jacob_carlson@g.harvard.edu

good in a mean square error sense), it will bias downstream estimates of economic quantities of interest.

In light of such a concern, both [Johnston et al. \[2026\]](#) and [Hitzig et al. \[2026\]](#) investigate the performance of their classifiers with what one might call *benchmark data*.

Example 1 (Continued). [Johnston et al. \[2026\]](#) use a benchmark dataset of 1000 random Codeforces problems, which contains pairs of coding tasks and (fastest) completion times for competitive coders. They find a correlation of 0.65 between the model-predicted times and the average completion time among the five fastest human solvers of the coding task. $\triangle$

Example 2 (Continued). [Hitzig et al. \[2026\]](#) use a benchmark dataset of 999 freelance job postings selected via stratified sampling over various price tiers. They find a Pearson R^2 between log actual and log predicted price of 0.38 in the full stratified sample. $\triangle$

Example 3 (Continued). [Hitzig et al. \[2026\]](#) use a benchmark dataset (of undisclosed size) consisting of internal Anthropic Claude Code sessions, for which sessions have observed ground truth success outcomes (i.e., whether the session’s commits landed on the main branch). They find that classifier judged success rises by roughly 3–5 percentage points per expertise level (with author fixed effects), and sessions at higher classifier rated expertise remain more likely to produce code that is successful. Additionally, [Hitzig et al. \[2026\]](#) construct a benchmark dataset based on SWE-chat [[Baumann et al., 2026](#)], where ground truth is imputed via a “strong model” (Mythos Preview). There is 82% agreement between the strong model ground truth success label and the success classifier used in the main report. $\triangle$

These benchmark data analyses clearly provide us with some signal that the measurement error in the privacy-preserving classifiers isn’t catastrophic. However, it is unclear to what extent this observed measurement error in benchmarks translates into error in the actual estimates of economic parameters of interest: the regression coefficients and means that the papers compute downstream of these classifier predictions. And in a sense, this is what we really care about: unbiased, high confidence estimation of the economically meaningful parameters, not good classification for good classification’s sake, for some particular performance metric.

A natural question is then: might we be able to use these benchmark datasets to help correct for the estimation-relevant consequences of measurement error in these privacy-preserving classifiers? The intuition for how we might do this is somewhat straightforward: we can get some sense of how systematically wrong the classifier is on the benchmark dataset, and perhaps use that knowledge of wrongness to adjust our estimator computed from/for the target AI usage dataset—so long as we take into account that the types of errors made in the benchmark dataset may not be totally representative of those in the usage dataset due to *distribution shift*. The purpose of this note will be to show that we can do exactly this wrt almost any downstream estimation goal of economic interest (for certain benchmark datasets, under certain assumptions).

Before moving on to specifics, it’s worth commenting on the fact that an alternative idea here would be to subsample from the AI usage dataset itself, and invoke some expensive

process to label that subsample with ground truth values—which some literatures would call a *validation dataset*. That way, we ensure the classifier error we are measuring is representative in the way we want. Indeed, when one can do this, there are many statistical benefits to using a validation sample to correct for measurement error in downstream estimation [Angelopoulos et al., 2023, Egami et al., 2024, Carlson and Dell, 2025]. However, specifically in this setting where *privacy-preservation* is required, and no human should be looking at real AI usage data, this route may not be feasible. Still, we might sometimes simply find the ground truth signal in a given benchmark to be most meaningful, e.g., humans carefully guessing how long it would take to complete a task versus actual task completion timestamps, or Mythos Preview’s judgment of success versus an actual commit signal from internal lab data. That said, to the extent that automated, privacy-preserving means exist to collect meaningful ground truth for a subset of the target AI usage dataset, such a strategy could be preferred. (In the above settings, it is only Example 3 that would permit doing so, using Mythos Preview as a ground truth signal for a small portion of the full session dataset.)

2 Warm-up

Let’s add some notation here. Say that we have a dataset of unique AI users (for simplicity) indexed by i and text descriptions of their task pulled from a session transcript, X_i , at a given snapshot in time. Say these users are sampled iid. In principle, there exists some meaningful ground truth measurement Y_i^* for each unique user i , e.g., the experienced human task completion time associated with the task described in X_i , or the successful achievement of a task described in X_i . The problem is that we don’t get to see Y_i^* for any of the individuals in the AI usage dataset. This is a *missing data* problem.

However, we do have access to a benchmark dataset with some ground truth, i.e., we have some (Y_i^*, X_i) pairs, just drawn from a different distribution—the distribution from which the benchmark dataset was drawn. Let’s join our target user and benchmark datasets, and for record keeping purposes we’ll introduce the random variable D_i , which indicates whether that observation i is from the AI usage dataset of interest ($D_i = 1$) or from the benchmark dataset ($D_i = 0$). With this combined dataset convention, we can write the estimand (parameter) of interest—say, for Example 1, the average human task completion time in the target AI user population—as

$$\theta := E[Y^* \mid D = 1].$$

Recall we only get to see Y^* when $D = 0$ under this bookkeeping convention. To formally encode this missingness, we write that we get to see the random variable Y_i for *every* i in the combined dataset, where

$$Y_i = (1 - D_i)Y_i^*.$$

That is, when $D_i = 0$ the $Y_i = Y_i^*$, but otherwise when $D_i = 1$ the $Y_i = 0$, where zero is just a convention for a missing value. (Note we can always differentiate a meaningful zero from a missingness zero by looking at D_i .)

As such, θ is not “identified” under our missing dataset setup (i.e., it is not computable using our observed data, even if we had as much of it as we want) *unless*, for instance, we

make an assumption that looks like the following:

$$[Y^* \perp\!\!\!\perp D] \mid X.$$

This resembles an “unconfoundedness” assumption in the causal inference literature (another quintessential missing data problem) but can also be seen substantively as a *covariate shift* assumption in the present context, which is that for all possible values any X_i can take on, $x \in \mathcal{X}$,

$$Y^* \mid (X = x, D = 1) \stackrel{L}{=} Y^* \mid (X = x, D = 0),$$

where $\stackrel{L}{=}$ denotes equality in law (distribution).¹ That is, the conditional distribution of $Y^* \mid X$ is “stable” across the AI usage and benchmark datasets, though the marginal distributions of X are able to vary. Concretely, this would mean that, for Example 1, conditional on the coding task, the distribution of human completion times is the same in the target and benchmark datasets, but the marginal distribution of coding tasks may differ across the target and benchmark datasets.

If we can reasonably assume this condition, then we can leverage the stability of the relationship $Y^* \mid X$ across the target and benchmark distributions in order to debias our estimates of θ wrt systematic measurement error in the session classifier. Doing this well will come down to how well we can effectively understand the covariate shift, the change in the marginal distribution of X . Fortunately, in practice, this will just mean we need to build one more classifier, this time of whether a given task belongs in the target AI usage or benchmark dataset. Finally, note that what we can debias effectively then is a Y^* that is conditionally stable, e.g., for Example 1, completion times for competition coders, or for Example 2, the price of the task-as-posting in a particular freelancer market. We’ll return to this observation in the final section of the note.

3 Framework

Now that we’ve established some basic notation and intuition for the strategy of handling systematic measurement error, we turn to an overview of a fully general and rigorous framework for accomplishing our goals.

In particular, we will adapt the framework of [Chen et al. \[2008\]](#), in the spirit of [Carlson and Dell \[2025\]](#). Specifically, we adapt the “verify-out-of-sample” results of [Chen et al. \[2008\]](#). These results will actually allow us to correct for measurement error in privacy-preserving classifier predictions used to estimate a parameter defined by *any* GMM identification strategy (so almost anything of interest: means, linear regression/model coefficients, structural model parameters, and more).

To start, we use the same essential setup as [Chen et al. \[2008\]](#), generalizing what we did in our warm-up. We start with an iid sample of data of n sessions indexed by i , $\{(X_i, Y_i, Z_i, D_i)\}_{i=1}^n$, where X_i is some appropriate excerpt or summary of AI session text data, Z_i are other covariates for the observation (say, other structured metadata about the session), and D_i keeps track of where the observation came from: $D_i = 1$ if the observation

¹This is an if and only if relationship (see Proposition 1 in the appendix).

came from the target dataset, and $D_i = 0$ if it came from a benchmark dataset. We again define $Y_i = (1 - D_i)Y_i^*$. To simplify some notation later on, we'll also introduce the sets $\mathcal{A} := \{i : D_i = 1\}$ and $\mathcal{B} := \{i : D_i = 0\}$, the indices of observations in the target and benchmark datasets, respectively.

Say we are interested in estimating a parameter $\theta \in \Theta \subseteq \mathbb{R}^{d_\theta}$ given by the moment condition

$$E[m((Y^*, Z); \vartheta) \mid D = 1] = 0 \quad \text{iff} \quad \vartheta = \theta$$

where $m(\cdot; \vartheta)$ is a set of functions with dimension $d_m \geq d_\theta$.

Going forward, let's label $\tilde{X} := (X, Z)$. Then the more general assumption we make to identify θ from observed data is just

$$[Y^* \perp\!\!\!\perp D] \mid \tilde{X}.$$

Under this assumption, by the law of total expectation,

$$\begin{aligned} E[m((Y^*, Z); \vartheta) \mid D = 1] &= E[E[m((Y^*, Z); \vartheta) \mid D = 1, (X, Z)] \mid D = 1] \\ &= E[E[m((Y^*, Z); \vartheta) \mid D = 0, \tilde{X}] \mid D = 1] \\ &= E[E[m((Y, Z); \vartheta) \mid D = 0, \tilde{X}] \mid D = 1], \end{aligned}$$

which is a moment condition solely in terms of the observed data. We also need positivity assumptions here so that this expectation exists, sufficiently that

$$p := P(D = 1) \in (0, 1), \quad p(\tilde{X}) := P(D = 1 \mid \tilde{X}) \in (0, 1) \text{ a.s.}$$

Because $\tilde{X}$ is plausibly quite high-dimensional (it is text data from an AI session), this condition—sometimes called “overlap”—may plausibly fail to hold [D’Amour et al., 2021], or fail in practice (causing high variance/finite sample instability). That said, the condition suggests that we want to *choose* a benchmark dataset that has high covariate overlap (session similarity) with our target usage dataset, all else equal.² Otherwise, we want to deploy strategies at our disposal (mostly borrowed from the causal inference literature) to ameliorate this problem, e.g., change the parameter to one that is conditioned on a high(er) overlap region of the covariate space; “trim” the $p(\tilde{X})$; or invoke representation learning strategies to lower the dimensionality of $\tilde{X}$ without introducing more bias [Clivio et al., 2026].³

We know from Chen et al. [2008] what the “efficient influence function” (EIF) Ψ for any θ in this framework is:

$$\Psi(Y, \tilde{X}, D; \theta) = -\mathcal{J}_\theta^{-1} \times F_\theta(Y, \tilde{X}, D)$$

where

$$F_\theta(Y, \tilde{X}, D) = \frac{1 - D}{p} \frac{p(\tilde{X})}{1 - p(\tilde{X})} [m((Y, Z); \theta) - \mathcal{E}(\tilde{X}; \theta)] + \frac{\mathcal{E}(\tilde{X}; \theta)}{p} D$$

for

$$\mathcal{E}(\tilde{X}; \theta) = E[m((Y^*, Z); \theta) \mid \tilde{X}], \quad \mathcal{J}_\theta = \frac{\partial}{\partial \theta} E[m((Y^*, Z); \theta) \mid D = 1].$$

²Technically, a minimal viable condition here is just that $P_{\tilde{X}|D=1} \ll P_{\tilde{X}|D=0}$. That is, we should choose a benchmark that *covers* the target tasks, but does not necessarily need to *match* them.

³At a minimum, this condition suggests that we want to push the benchmark and target dataset X_i through the same formatter/summarizer.

What is this EIF object, and why do we care about it? We care about the EIF in this setting primarily because it can be used to derive a “doubly robust,” Neyman orthogonal estimator [Chen et al., 2026], which is permissive of using slightly biased (but overall high quality) black-box AI/ML models to estimate the “nuisance” functions that appear in our empirical setting, i.e., the $p(\tilde{X}), \mathcal{E}(\tilde{X}; \theta)$ here.⁴ Concretely, it is also true that

$$E[\Psi(Y, \tilde{X}, D; \vartheta)] = 0 \quad \text{iff} \quad \vartheta = \theta,$$

and so Ψ can play the same role as m , but the estimating equation it suggests will yield a more robust estimator. Moreover, estimators derived from EIF-based estimating equations will typically have the lowest possible variance of any nice, “regular” estimator of our parameter of interest. As such, there is a very real sense that, if we want to use benchmark data to debias our estimator of a quantity of interest θ , we can do no better from a robustness and efficiency standpoint than an estimator derived using the above framework.

4 Applying the Framework

With a general framework in place, let’s go revisit some of our motivating examples, and consider estimation objectives in them which can be cast as estimating a mean of missing ground truth data.

Example 1 (Continued). We want to estimate the expected human completion time for coding tasks in some period of AI use. Here, $Y^* \in \mathbb{R}$ is a human completion time, and thus the parameter of interest is $\theta = E[Y^* \mid D = 1]$. $\triangle$

Example 2 (Continued). We want to estimate the average price of tasks encoded in AI sessions for some period of AI use. Letting $Y^* \in \mathbb{R}$ be the freelance market price of a task in the AI usage data, the parameter of interest is $\theta = E[Y^* \mid D = 1]$. $\triangle$

Example 3 (Continued). We want to estimate the probability of a successful AI coding session in some period of AI use. Here the $Y^* \in \{0, 1\}$ is a session success indicator, and the parameter of interest is $\theta = P(Y^* = 1 \mid D = 1) = E[Y^* \mid D = 1]$. $\triangle$

In all of these examples, the relevant moment condition is simply

$$m(Y^*; \vartheta) = Y^* - \vartheta$$

which implies that, using our covariate shift assumption,

$$\mathcal{E}(X; \theta) = E[Y^* \mid X] - \theta = E[Y \mid D = 0, X] - \theta, \quad \mathcal{J}_\theta = -1,$$

⁴These are “nuisance” functions in the sense that we don’t care about learning them for their own sake, per se—we just need to learn them in order to learn θ .

as here $\tilde{X} = X$. Thus we can compute that in this setting

$$\begin{aligned} F_\theta(Y, \tilde{X}, D) &= \frac{1-D}{p} \frac{p(X)}{1-p(X)} [m(Y; \theta) - \mathcal{E}(X; \theta)] + \frac{\mathcal{E}(X; \theta)}{p} D \\ &= \frac{1-D}{p} \frac{p(X)}{1-p(X)} [(Y - \theta) - (E[Y | D = 0, X] - \theta)] + \frac{E[Y | D = 0, X] - \theta}{p} D \\ &= \frac{1-D}{p} \frac{p(X)}{1-p(X)} (Y - E[Y | D = 0, X]) + \frac{E[Y | D = 0, X] - \theta}{p} D \end{aligned}$$

and so the EIF is

$$\Psi(Y, \tilde{X}, D; \theta) = \frac{1-D}{p} \frac{p(X)}{1-p(X)} (Y - E[Y | D = 0, X]) + \frac{D}{p} E[Y | D = 0, X] - \frac{D}{p} \theta.$$

Using the EIF to inform an estimating equation immediately suggests the estimator (based on only a few lines of algebra⁵):

$$\hat{\theta} = \left(\frac{1}{n} \sum_i \frac{D_i}{\hat{p}} \right)^{-1} \frac{1}{n} \sum_{i=1}^n \left(\frac{D_i}{\hat{p}} \text{Cls}(X_i) + \frac{1-D_i}{\hat{p}} \frac{\hat{p}(X_i)}{1-\hat{p}(X_i)} (Y_i - \text{Cls}(X_i)) \right),$$

where $\hat{p} = \frac{1}{n} \sum_i D_i$, and $\text{Cls}(X_i) = \hat{E}[Y | D = 0, X_i]$ and $\hat{p}(X_i) = \hat{P}(D = 1 | X_i)$, our best feasible approximations of $E[Y | D = 0, X_i]$, $P(D = 1 | X_i)$, respectively.

Simplifying,

$$\hat{\theta} = \frac{1}{n} \sum_{i=1}^n \left(\frac{D_i}{\hat{p}} \text{Cls}(X_i) + \frac{1-D_i}{\hat{p}} \frac{\hat{p}(X_i)}{1-\hat{p}(X_i)} (Y_i - \text{Cls}(X_i)) \right)$$

or equivalently, writing out terms as averages from the target and benchmark datasets, respectively,

$$\hat{\theta} = \frac{1}{|\mathcal{A}|} \sum_{i \in \mathcal{A}} \text{Cls}(X_i) + \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \frac{|\mathcal{B}|}{|\mathcal{A}|} \frac{\hat{p}(X_i)}{1-\hat{p}(X_i)} \{Y_i - \text{Cls}(X_i)\}.$$

Notice that the role of $\text{Cls}(x)$ is usually played by the privacy-preserving classifier, which in the status quo is implemented with a frontier LLM; the efficiency theory here tells us what the classifier should approximate is a particular conditional expectation function, $x \mapsto E[Y^* | X = x]$.⁶ The $\hat{p}(x)$ is another classifier we should build based on an independent hold-out sample—the probability mass from a “discriminating” binary classifier of the task arising in the target data versus the benchmark data (perhaps approximated using the softmax distribution over the final token of a calibrated LLM, or some small neural network trained on an independently held-out dataset). That is, there is *one* additional prediction problem we must do for appropriate handling of measurement error, which we can still plausibly do

⁵Mechanically, one takes the moment condition based on the EIF and replaces population quantities with sample quantities, solving out for the $\hat{\theta}$ that sets the equation equal to zero.

⁶As such, we may want to encode this notion directly into the prompt, i.e., “Predict the *expected* completion time...”

in privacy-preserving way. Further notice that the issue of poor “overlap” directly manifests in the properties of the weights $\frac{\hat{p}(x)}{1-\hat{p}(x)}$.

For comparison, the (possibly biased) estimator implicitly being used in the status quo, which we might call $\tilde{\theta}$, is just the first term in $\hat{\theta}$, the

$$\tilde{\theta} := \frac{1}{|\mathcal{A}|} \sum_{i \in \mathcal{A}} \text{Cls}(X_i).$$

The work above informs us of the appropriate debiasing correction for $\tilde{\theta}$ using some benchmark data—the term in blue—assuming that we can model well the probability a given task comes from the benchmark or comes from the target AI usage dataset.

Recall, this machinery is general to any θ of economic interest defined by a GMM identification condition, though we have fully worked out the case for a simple mean estimation question. For instance, in the setting of Example 3, the θ of interest as implemented in [Hitzig et al. \[2026\]](#) is a regression coefficient; this could also be readily handled by the framework (though we omit the derivation for brevity of the note).

Also for the sake of brevity of this note, we will not derive the large-sample properties of this estimator, but proving a couple formal inference statements would be a straightforward next step. That said, this is almost certainly the actual estimator one would implement, given how it was derived—there are simply some additional questions about how well one needs to approximate $E[Y^* | X = x]$ and $p(x)$ that formal asymptotic analysis would clarify. If we knew $E[Y^* | X = x]$ and $p(x)$ with certainty, the inference results here would be immediate.

5 More Intuition

For more intuition, we can observe that the estimator $\hat{\theta}$ also naturally corresponds to a sensible estimator that measures the systematic mistakes in the benchmark predictions and then applies that as a bias correction to predictions on the target dataset, appropriately adjusting for covariate shift across benchmark and target data with a density ratio (Radon-Nikodym derivative).

To see this, consider another identification for θ the mean of missing data. In particular, say P is the target AI usage distribution of interest, and Q is the distribution of the benchmark data. We make a covariate shift assumption as before, now denoting it as $P_{Y|X} = Q_{Y|X}$ but where P_X and Q_X are able to differ. Then we have that

$$\theta := E_P[Y^*] = \underbrace{E_P[\text{Cls}(X)]}_{\text{Expected classifier prediction}} + \underbrace{E_Q[w(X)(Y^* - \text{Cls}(X))]}_{\text{Bias in benchmark, reweighted for cov. shift}}$$

where $w(x) = \frac{dP_X}{dQ_X}(x)$, the Radon-Nikodym derivative.

The natural plug-in estimator implied here, call it $\check{\theta}$, is just (assuming we knew the density ratio and p for now)

$$\check{\theta} = \frac{1}{n} \sum_{i=1}^n \left(\frac{D_i}{p} \text{Cls}(X_i) + \frac{1 - D_i}{1 - p} w(X_i)(Y_i^* - \text{Cls}(X_i)) \right).$$

By Bayes, however,

$$w(x) = \frac{dP_X}{dQ_X}(x) = \frac{f(x \mid D = 1)}{f(x \mid D = 0)} = \frac{P(D = 1 \mid x)f(x)/p}{P(D = 0 \mid x)f(x)/(1-p)} = \frac{p(x)(1-p)}{(1-p(x))p}$$

and so if we estimate p with $\hat{p} = n^{-1} \sum_i D_i$ and $p(x)$ with $\hat{p}(x)$

$$\check{\theta} = \frac{1}{n} \sum_{i=1}^n \left(\frac{D_i}{\hat{p}} \text{C1s}(X_i) + \frac{1-D_i}{\hat{p}} \frac{\hat{p}(X_i)}{1-\hat{p}(X_i)} (Y_i^* - \text{C1s}(X_i)) \right).$$

Notice that $\hat{\theta} = \check{\theta}$; it is equivalent to the estimator we found earlier with the verify-out-of-sample approach! Thus the intuition that allowed us to build $\check{\theta}$ is the right intuition for $\hat{\theta}$ and the verify-out-of-sample approach too.

6 Takeaways

What are the overall takeaways from this note?

1. To estimate something like the average complexity of tasks for some target AI user population, researchers are currently using a privacy-preserving automated classifier in an estimator that looks like:

$$\tilde{\theta} := \frac{1}{|\mathcal{A}|} \sum_{i \in \mathcal{A}} \text{C1s}(X_i).$$

This estimator may suffer from bias due to systematic measurement error in C1s . To purge this bias, we could leverage a benchmark dataset, so long as it only differs from the target population of interest by a covariate shift (change in marginal distribution of the tasks in the session); a highly general framework based on this assumption would tell us to use the debiased (and efficient) estimator:

$$\hat{\theta} = \frac{1}{|\mathcal{A}|} \sum_{i \in \mathcal{A}} \text{C1s}(X_i) + \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \frac{\hat{p}(X_i)}{|\mathcal{A}|(1-\hat{p}(X_i))} \{Y_i - \text{C1s}(X_i)\}.$$

Naturally, if there is no measurement error, $Y_i - \text{C1s}(X_i) = 0$ almost surely, and we recover the unadjusted estimator.

2. The theory that underlies this general framework would make it easy to debias any θ of interest fairly straightforwardly, be it a θ from a linear model, structural model, etc.
3. The theory underlying the framework reveals that we should ideally be choosing a benchmark dataset that has transcripts/tasks that are as similar as possible (in a distributional sense) to the ones in the target population.
4. The framework tells us we only need to do one more piece of engineering to debias our estimates: fit a discriminating classifier of whether an observation came from the target or benchmark dataset. Plausibly, this can be done in a privacy-preserving way.

5. The benchmark outcome (the “ground truth” the benchmark contains) naturally dictates which attribute of the target AI usage dataset we can most credibly leverage for estimation, e.g., per Example 1, if the outcome in the benchmark is competition coder completion times, that’s the outcome we can most credibly debias wrt (as opposed to completion times for average-skill coders), and the choice of privacy-preserving classifier prompt should reflect this. By the same logic, then, in the setting of Example 1, curating a dataset of real-world reminiscent coding tasks and timing their completion by non-competitive coders could act as a preferable benchmark.
6. The core covariate shift assumption we make in this framework is infringed when the benchmark is selected on Y^* , as this violates the comparability of $Y^* \mid X$ across samples. Selecting on (functions of) X , however, is benign. Thus, when it comes to Example 2, this suggests that we might want to stratify on the price *predicted* from X instead of using Y^* directly.
7. If we could annotate at least some of our AI sessions of interest with meaningful ground truth in privacy-preserving way, that would be the most promising thing to do, and the framework for doing so would look nearly identical to the one discussed in this note.⁷ However, the present approach provides a naturally privacy-preserving alternative, involving no human looking at real AI session data.

⁷We would essentially just use the “verify-in-sample” results from Chen et al. [2008], as documented in detail in a forthcoming version of Carlson and Dell [2025].

7 Appendix

Proposition 1. Assume that $P(D = 1) \in (0, 1)$, $P(D = 1 \mid X) \in (0, 1)$ a.s. Then $[Y^* \perp\!\!\!\perp D] \mid X$ iff $Y^* \mid (X, D = 1) \stackrel{L}{=} Y^* \mid (X, D = 0)$.

Proof. Our goal is to show that almost surely

$$[Y^* \perp\!\!\!\perp D] \mid X \iff Y^* \mid (X, D = 1) \stackrel{L}{=} Y^* \mid (X, D = 0).$$

The first direction then is

$$[Y^* \perp\!\!\!\perp D] \mid X \implies Y^* \mid (X, D = 1) \stackrel{L}{=} Y^* \mid (X, D = 0).$$

To show this, it is immediate that assuming $[Y^* \perp\!\!\!\perp D] \mid X$ implies

$$P(Y^* = y \mid X, D = 1) = P(Y^* = y \mid X) = P(Y^* = y \mid X, D = 0).$$

To show the other direction,

$$[Y^* \perp\!\!\!\perp D] \mid X \iff Y^* \mid (X, D = 1) \stackrel{L}{=} Y^* \mid (X, D = 0)$$

observe that if $Y^* \mid (X, D = 1) \stackrel{L}{=} Y^* \mid (X, D = 0)$, then

$$P(Y^* = y, D = d \mid X) = P(Y^* = y \mid X, D = d)P(D = d \mid X)$$

and then that by the law of total probability

$$P(Y^* = y \mid X) = P(Y^* = y \mid X, D = 1)P(D = 1 \mid X) + P(Y^* = y \mid X, D = 0)P(D = 0 \mid X).$$

But under the covariate shift assumption, then for all $d \in \{0, 1\}$

$$\begin{aligned} P(Y^* = y \mid X) &= P(Y^* = y \mid X, D = 1)P(D = 1 \mid X) + P(Y^* = y \mid X, D = 0)P(D = 0 \mid X) \\ &= P(Y^* = y \mid X, D = d)P(D = 1 \mid X) + P(Y^* = y \mid X, D = d)P(D = 0 \mid X) \\ &= P(Y^* = y \mid X, D = d)(P(D = 1 \mid X) + P(D = 0 \mid X)) \\ &= P(Y^* = y \mid X, D = d) \end{aligned}$$

and so

$$\begin{aligned} P(Y^* = y, D = d \mid X) &= P(Y^* = y \mid X, D = d)P(D = d \mid X) \\ &= P(Y^* = y \mid X)P(D = d \mid X) \end{aligned}$$

which is the definition of conditional independence, $Y^* \perp\!\!\!\perp D \mid X$. □

References

- Drew Johnston, David Holtz, Alex Martin Richmond, Christopher Ong, Prasanna Tambe, and Aaron Chatterji. The Shift to Agentic AI: Evidence from Codex, June 2026. URL <https://arxiv.org/abs/2606.26959>. Version Number: 1.
- Zoe Hitzig, Maxim Massenkoff, Eva Lyubich, Ryan Heller, and Peter McCrory. Agentic coding and persistent returns to expertise, 2026. URL <https://www.anthropic.com/research/research/claude-codeexpertise>.
- Joachim Baumann, Vishakh Padmakumar, Xiang Li, John Yang, Diyi Yang, and Sanmi Koyejo. Swe-chat: Real-world ai coding sessions in the wild. *arXiv preprint arXiv:2604.20779*, 2026. URL <https://arxiv.org/pdf/2604.20779>.
- Anastasios N. Angelopoulos, Stephen Bates, Clara Fannjiang, Michael I. Jordan, and Tijana Zrnic. Prediction-powered inference. *Science*, 382(6671):669–674, November 2023. doi: 10.1126/science.adi6000. URL <https://www.science.org/doi/10.1126/science.adi6000>.
- Naoki Egami, Musashi Hinck, Brandon M. Stewart, and Hanying Wei. Using Large Language Model Annotations for the Social Sciences: A General Framework of Using Predicted Variables in Downstream Analyses. June 2024. URL https://naokiegami.com/paper/dsl_ss.pdf.
- Jacob Carlson and Melissa Dell. A Unifying Framework for Robust and Efficient Inference with Unstructured Data, 2025. URL <https://arxiv.org/abs/2505.00282>. Version Number: 2.
- Xiaohong Chen, Han Hong, and Alessandro Tarozi. Semiparametric efficiency in GMM models with auxiliary data. *The Annals of Statistics*, 36(2), April 2008. ISSN 0090-5364. doi: 10.1214/009053607000000947. URL <https://projecteuclid.org/journals/annals-of-statistics/volume-36/issue-2/Semiparametric-efficiency-in-GMM-models-with-auxiliary-data/10.1214/009053607000000947.full>.
- Alexander D’Amour, Peng Ding, Avi Feller, Lihua Lei, and Jasjeet Sekhon. Overlap in observational studies with high-dimensional covariates. *Journal of Econometrics*, 221(2): 644–654, April 2021. ISSN 0304-4076. doi: 10.1016/j.jeconom.2019.10.014. URL <https://linkinghub.elsevier.com/retrieve/pii/S0304407620302694>.
- Oscar Clivio, Alexander D’Amour, Alexander Franks, David Bruns-Smith, Chris Holmes, and Avi Feller. Deconfounding Scores and Representation Learning for Causal Effect Estimation with Weak Overlap, 2026. URL <https://arxiv.org/abs/2604.00811>. Version Number: 1.
- Yuxi Chen, Edward H. Kennedy, and Sivaraman Balakrishnan. On the Equivalence between Neyman Orthogonality and Pathwise Differentiability, 2026. URL <https://arxiv.org/abs/2603.15817>. Version Number: 1.